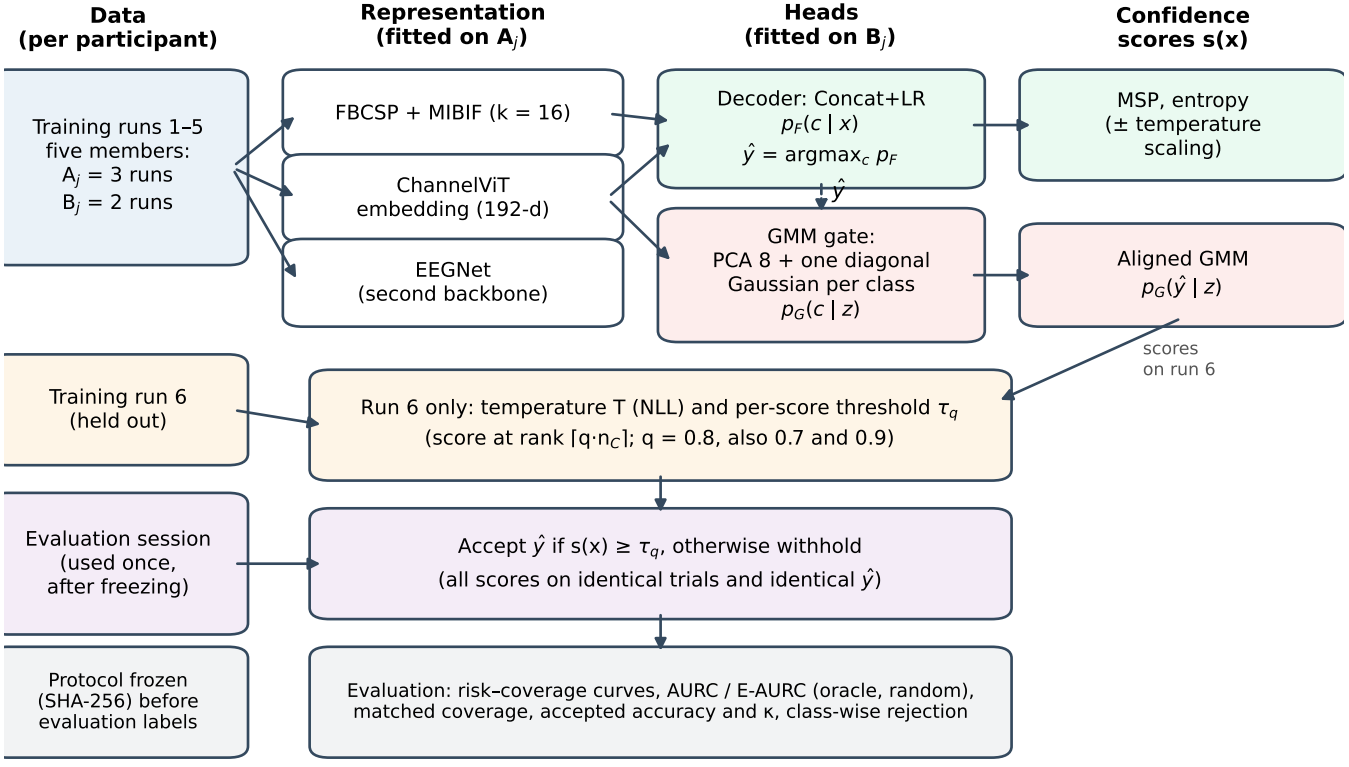
Data
(per participant)
Representation
(fitted on A_j)
Heads
(fitted on B_j)
Confidence
scores s(x)
Training runs 1–5
five members:
A_j = 3 runs
B_j = 2 runs
FBCSP + MIBIF (k = 16)
ChannelViT
embedding (192-d)
EEGNet
(second backbone)
Decoder: Concat+LR
$p_F(c \mid x)$
$\hat{y} = \text{argmax}_c\ p_F$
$\hat{y}$
GMM gate:
PCA 8 + one diagonal
Gaussian per class
$p_G(c \mid z)$
MSP, entropy
(± temperature
scaling)
Aligned GMM
$p_G(\hat{y} \mid z)$
scores
on run 6
Training run 6
(held out)
Run 6 only: temperature T (NLL) and per-score threshold τ_q
(score at rank $\lceil q \cdot n_C \rceil$; q = 0.8, also 0.7 and 0.9)
Evaluation session
(used once,
after freezing)
Accept $\hat{y}$ if s(x) ≥ τ_q, otherwise withhold
(all scores on identical trials and identical $\hat{y}$)
Protocol frozen
(SHA-256) before
evaluation labels
Evaluation: risk–coverage curves, AURC / E-AURC (oracle, random),
matched coverage, accepted accuracy and κ, class-wise rejection